

1 Introduction

Facebook, at the time of writing, is the world's most popular social network service with over 500 million users active in the last 30 days [10]. Encrypted Facebook is a Firefox extension which aims to protect privacy through the use of a broadcast encryption scheme when sharing certain content via the Facebook platform.

1.1 Background

Since its inception, Facebook has come under criticism in relation to on-line privacy [3]. As social networks mimic real life interactions members are often willing to reveal more private details than they otherwise would, resulting in social networks accumulating a large repository of sensitive information [8]. There are a number of ways this information can then be exposed: users misconfiguring privacy settings,¹ malice, error or neglect on the part of the social network and its employees,² acquiescent disclosure to government authorities and unlawful access through phishing scams and security exploits [4] [fb-gov] [11] [15].

Encrypting content ensures its secrecy in any eventuality, provided the encryption key itself is kept safe. Tools exist for performing encryption manually - some are even partly integrated into the Facebook user interface [firegpg]. However, studies have suggested that software which requires manual management of cryptographic keys is not usable enough to provide effective security for most users [johnny]. In addition, interaction over social networks is typically one-to-many, unlike applications like

¹Facebook was recently forced to update its privacy controls due to growing pressure from the public and media [12].

²As was the case in 2010 when personally identifiable Facebook user details were sold to data brokers by third party developers [5].

e-mail which these tools are designed for.

Solutions targetting social networks specifically have been proposed. Some are still overly reliant on manual key management; others fail to fully protect the user's privacy or support the most popular forms of content (see sections 1.4 and ??). Thus far all proposals have required the use of a third party in ways which are unlikely to be scaleable (see section xxx).

Attempts have been made to create new social networks which protect privacy by encrypting content and in some cases even decentralizing hosting of the entire platform [**pidder**] [**diaspora**]. Unfortunately, network externalities make it difficult for newcomers to compete with Facebook since the utility of a social network is intrinsically tied to the size of its userbase.³

1.2 Objectives

The aim of this project is to provide an enhanced privacy solution for existing Facebook users which is:

- **Privacy preserving** given the scenarios presented in section 1.1.
- **Useable** and therefore accessible to the typical Facebook user.
- **Scaleable** given the vast size of Facebook's userbase.
- **Incrementally deployable** in order to avoid some of the problems of network effects.

1.3 Limitations

Some objectives we consider non-goals, or outside the scope of the project:

- Securing the actual structure of the social graph.
- Ensuring the integrity, authenticity or non-repudiation of communications. In theory the public key scheme could be extended to do so but this would go beyond mere privacy control.

³Some suggest the value of a social network grows linearithmically, quadratically or even exponentially with the number of users [14] [9].

- Protection against middleperson attacks (see section XXX).
- Guaranteeing availability of content. Completely severing reliance on Facebook would be tantamount to building a new network.
- Concealing the existence of content itself or concealing the fact that it is encrypted. Steganography is employed but not for this reason (see section XXX).
- Full cross-platform and/or cross-browser support (see section XXX).
- Designing and implementing a security policy which comprehensively covers all aspects of the project's functionality (see section XXX).

1.4 Existing work

Many applications for encrypting online data exist, most in the form of browser extensions [**firepgg**] [**cryptfire**] [**textcrypt**]. There are also several applications which target Facebook specifically.

uProtect.it Client side JavaScript application which inserts UI controls and intercepts content. Content is encrypted, stored and decrypted all on a third party server.
<http://uprotect.it>

FaceCloak Firefox extension which posts fake content to Facebook, using it as an index to the encrypted content on a third party server. Running your own server is possible, though only users on the same server can communicate.
<http://crysp.uwaterloo.ca/software/facecloak/>

flyByNight Content is submitted through a Facebook application, encrypted using client side JavaScript and passed via Facebook to a third party application server. Proxy re-encryption is used for sending to multiple recipients.
<http://hatswitch.org/~nikita/>

NOYB Content is stored in plaintext but profiles are anonymised. The mapping from real-to-fake profiles is known only to a user's friends.
<http://adresearch.mpi-sws.org/noyb.html>

2 Preparation

In this chapter we formulate the objectives presented in section ?? into a set of design principles, drawing on existing work. We justify the use of broadcast encryption and define a suitable scheme. We describe possible deployment strategies and briefly review Firefox extension development and the Facebook platform. We also look at the specific problems associated with encrypting images and possible solutions. Finally, we discuss some additional security related concerns and draft a threat model, outline an appropriate testing strategy and software development methodology and derive a concrete set of requirements.

2.1 Design principles

We describe the design principles of Encrypted Facebook; how they enforce the stated aims of privacy preservation, scalability, usability and incremental deployment and how they compare with existing approaches. Note that here use of the term third party excludes Facebook itself.

2.1.1 Encryption of shared content

It is possible to preserve privacy by encrypting or otherwise concealing the link between a real life user and their online identity, as with NOYB [**noyb**]. Arguably, however, privacy is only poorly preserved due to problems of inference control [**ross**]. An example: many users will be easily identifiable simply from the photos they upload. Incremental deployment is also not possible since non-users will only ever see fictitious profiles [**facecloak**].

The alternative is to encrypt shared content itself in some way or another, restricting access to only those who possess the appropriate key even in the event of release.

2.1.2 Independency from third-party servers

In addition to encrypting content, flyByNight, FaceCloack and uProtect.it all opt to migrate content from the Facebook platform to an external third party database.

With Encrypted Facebook we avoid outsourcing content due to scalability concerns. Storing and delivering encrypted content requires at least the resources needed for storing and delivering the cleartext. Facebook's monthly bandwidth overheads alone are in excess of two million dollars [**fb-costs**]. They are able to offer a free service by serving highly targetting advertising to members based on the structure of the social graph [**fb-ads**]. This revenue stream would be largely unavaible to any solution hosting a database of encrypted content. A subscription based service would also be unlikely to scale, since the majority of Facebook users would be reluctant to pay [**fb-pay**].

Third-party servers can also be employed for performing encryption and/or decryption, as with uProtect.it and flyByNight. We avoid taking this approach, again since the resources required by the server would scale linearly in proportion to the amount of content exchanged.

Another use of third-party servers is to store and distribute cryptographic keys, perhaps as part of a public key infrastructure. Arguably issues of scale here are not so severe. In any case we do not rely on such a service due to conflicts with other design principles, or in some cases due to a limited percieved benefit given the additional complexity of im-
plementation.

2.1.3 Secret key security

Any encryption scheme will require some form of key whose secrecy is required.

It is possible to use a trusted third-party to store and distribute secret keys, in a so-called key escrow arrangement. Key management can even be taken out of the users hands entirely, improving useability. This is the basis of uProtect.it [**uprotect**]. However, confidentiality is only weakly assured: trust has simply been deferred from Facebook to the third-party and many of the scenarious raised in section 1.1 still apply.

Another possibility is using secret keys derived from a password. fly-ByNight, for example, allows users to download a password-protected private key from their server. We could also store password protected keys in-band (i.e. on Facebook itself) or generate a key simply by hashing a

password. Again, relying on the user to memorise a password rather than manage secret keys improves useability. Unfortunately the entropy of user chosen passwords is far less than that of randomly generated keys [password].

With Encrypted Facebook we ensure keys are only ever generated and stored on the user's device(s), trading useability for better privacy protection.

2.1.4 Minimal use of OOB channels

Secure **OOB!** (OOB!) channels (such as encrypted email or face-to-face exchange) can be used to transmit update messages, keys or other information as part of an encryption scheme (see section XXX). Since these channels are, by definition, external to the Facebook platform it can be hard to automate such exchanges and much is still required from the user. FaceCloak, for example, requires users to transmit messages over secure email when adding friends. The process is partly automated, however the user must set up an email client themselves and ensure that the email is sent securely (by using PGP, for example).

Encrypted Facebook is designed to limit the use of **OOB!** transmission due to useability concerns regarding manual key management. The exception is when installing a secret key across more than one device - since transporting a secret key by any other method would compromise the principle of secret key security.

2.2 Broadcast encryption

Communication over social networks is typically one-to-many whereas cryptography traditionally considers a sender and a single recipient. We look at existing solutions to this problem and outline the broadcast encryption scheme adopted by Encrypted Facebook and a justification for its use.

2.2.1 Existing solutions

Existing proposals tackle this problem in the following ways:

- If content is both hosted and encrypted/decrypted remotely, as with uProtect.it, one-to-many support is trivial. The user simply authenticates with the server and is sent the cleartext.

- If a third-party can be used for computation we can use a technique called proxy re-encryption as with [flybynight]. Here the server changes the key under which the content may be decrypted on demand, without ever being able to read the cleartext itself [proxy].
- Distributing keys over **OOB!** channels can permit one-to-many communication. A FaceCloak user, for example, shares a single decryption key **OOB!** among friends.

None of these approaches are compatible with Encrypted Facebook's design principles. Instead we use a form of broadcast encryption.

2.2.2 Proposed scheme

Appendix XXX gives a full introduction to broadcast encryption schemes. In general they can be characterised by the size of each users private and public key and the amount of transmission overhead that must be sent with each message, given the number of recipients. The scheme presented here has a large transmission overhead, however unlike many more complex schemes it doesn't require the use of key-update messages. Transmitting key-updates **OOB!** or though a third party would violate our design principles. Transmitting them in-band would be possible, but since Facebook is a "best effort" service synchronisation and lost updates would likely be a problem. Further discussion of more advanced schemes is included in appendix XXX.

Let U be the set of users with Encrypted Facebook installed and $R \subseteq U$ be a set of intended recipients. Our broadcast encryption scheme is defined as follows:

Definition 2.1. Given a suitable asymmetric encryption scheme P and a suitable symmetric scheme Q , we define our broadcast encryption scheme as the triple of algorithms (SETUP, BROADCAST, DECRYPT) such that:

- The setup algorithm (SETUP) takes a user $u \in U$ and constructs that receiver's private key $priv_u$ and public key pub_u using scheme P .
- The broadcast algorithm (BROADCAST) takes the list of privileged users R and a message m , generates a session key k using scheme Q and broadcasts a message $b = (b_1, b_2)$ where:
 - b_1 is the list of pairs (u, k_u) such that $u \in R$, where k_u is the session key encrypted under pub_u .

- b_2 is m encrypted under a session key k .
- A user $u \in U$ runs the decryption algorithm $\text{DECRYPT}(b, u, \text{priv}_u)$ that will:
 - If $(u, k_u) \in b_2$, extract the session key k from k_u using priv_u . m is obtained by decrypting b_2 using k
 - DECRYPT fails, if $(u, k_u) \notin b_2$ or if, equivalently, $u \notin R$.

2.3 Intercepting Facebook interactions

In order to encrypt content it must be intercepted before being submitted to Facebook. We describe the possible stages at which this may occur (Figure 2.1):

- a) On a remotely hosted proxy server. Easier multi-OS and multi-browser support.
- b) On a proxy server running on localhost. Easier multi-browser support.
- c) Within the browser, outside the browser sandbox. Extensions and plugins exist here and have elevated privileges over normal site code. Other examples include signed Java applets, ActiveX controls and to a lesser extent inline Flash and Silverlight applications.¹ FaceCloak takes this approach.
- d) Within the browser, entirely inside the browser sandbox - using only JavaScript and HTML. uProtect.it and flyByNight both take this approach.
- e) Outside the browser as part of a bespoke Facebook client application.

(a) conflicts with the design goal of no server-side computation. The browser sandbox prevents local filesystem access, ruling out (d) if private keys are to be kept securely. We take approach (c) since (b) and (e) are considerably more complex, at some cost to cross-platform compatibility.

¹Flash applications, for example, are restricted but can provide basic filesystem access [**flash-sbox**].

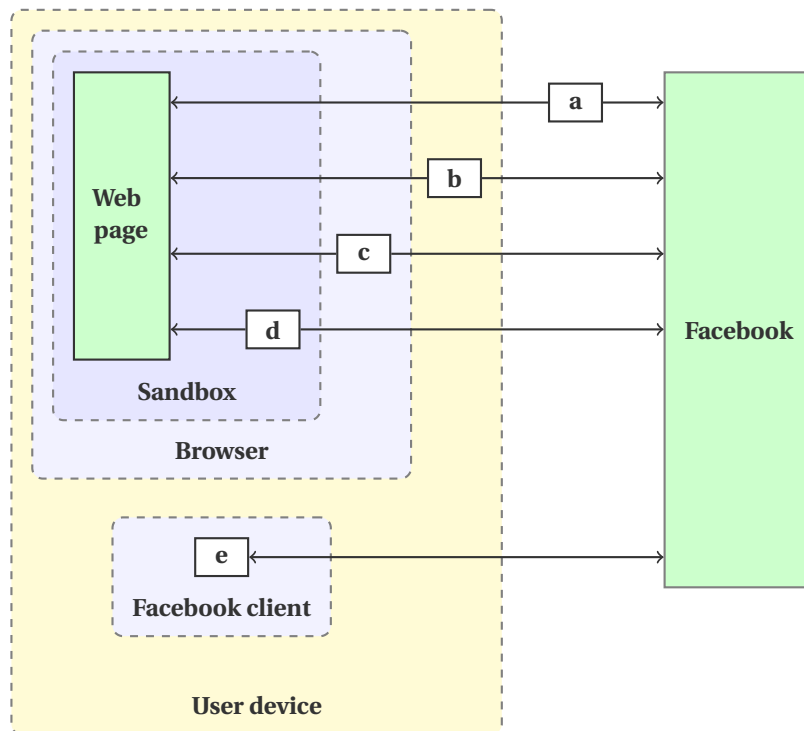


Figure 2.1: Possible deployment strategies for intercepting interaction with Facebook.

2.4 Mozilla Firefox extension development

Encrypted Facebook is developed as an extension for Mozilla Firefox, the worlds most popular browser (as of January 2011). Manipulating the **DOM! (DOM!)** is well supported by browser extensions since the browser interface chrome is often built on existing web technologies. This allows Encrypted Facebook to integrate functionality into Facebook's own user interface which mitigates useability issues concerning key management. Porting to other browsers is discussed in section XXX.

Firefox extensions are written in JavaScript with partial support for binding library code written in Python or C/C++ [**ffox-lang**]. Performing cryptography in JavaScript is possible but comes with severe performance difficulties [**flybynight**]. Table 2.2 compares approximate performance for each language (full details given in appendix xxx).

Since long delays would hamper useability Encrypted Facebook uses C/C++ for computation intensive operations. Native code can be executed

Language	Library	Time (ms)
Python 2.6.6	pycryptopp 0.5.17	1,220 ms
C++ 98	Botan 1.8.11	92 ms
JavaScript 1.6 (in Chrome 12.0.712)	JavaScript (last updated December 2005)	1,685,000 ms

Figure 2.2: *Approximate time for 256-bit AES encryption of 1000 1.5 MiB random messages.*

from within Firefox in at least three ways:

- Creating an XPCOM component. These are linked against a single Gecko ² version; supporting multiple versions is possible but non-trivial [**xpcom**].
- Loading native libraries with **js-ctypes**. Introduced in Gecko 2.0 [**js-ctypes**].
- Using **nsiProcess** to invoke an external stand-alone application. Capturing output can be difficult.

Since building an entire XCPOM component would be excessive and give little advantage by way of multi-version compatibility, the newly introduced **js-ctypes** module is used to load native code. This restricts use of Encrypted Facebook to Firefox version 4.0 and above. Gecko 2.0 does also provides slightly better support for working with the local filesystem and manipulating images in the **DOM!**.

2.5 The Facebook platform

Facebook represents the social graph as objects and connections between them. Objects include users, photos, messages and events. All objects are assigned a single unique Facebook ID, and all objects (except users themselves) are associated with the ID of their owner. The owners privacy settings (global and per object) will determine who may access an object.

²Gecko is the layout engine used by Firefox.

Objects will also have connections to other users and other objects. For example, users can create discussion threads by commenting on objects or tag other users as being present in a photograph.

There are many ways users can interact with Facebook. We will make the generalisation that interaction amounts to creating and retrieving content - i.e. objects in the social graph. Our goal then is to encrypt and decrypt the properties of certain objects - e.g. the body of a message object. We do not attempt to encrypt or otherwise conceal connections between objects, this is a non-goal as stated in section XXX.

We describe the most popular forms of content submitted and discuss issues relating to the connectedness of user nodes and signal-to-noise ratio in the network. We then briefly look at methods of interfacing with the Facebook platform.

2.5.1 Content types

Encrypting all possible types of content would be prohibitively complex. We therefore aim to encrypt only those most frequently used. From Table 2.1 we can clearly see that these are Comments, Messages, Images and Posts.

Our broadcast encryption scheme described in section XXX adds size overheads to any content that will be encrypted. Without relying on a third-party server all overheads must be stored on Facebook itself, in some form or another. Images and blog-style notes are obvious targets for storage utilisation due to their high capacity limits (see Table 2.1). In particular, notes can contain over 120 KiB of information since each character represents one 16-bit unicode code point. Images are subject to lossy compression which is discussed in section XXX.

Each user's profile has an "About Me" field with a character limit of XXX. With no other obvious capacitous attribute that can be easily queried from user ID, this is an obvious place to store a user's public key.

Activity	Frequency (per second)	Limitations	Notes
Comment	8,507	8,000 chars.	
Message	2,263	10,000 chars.	
Image	2,263	720 × 720 pixels	3-channel 8-bit colour. JPEG compressed (see section XXX).
Friend request	1,643		Social graph structure.
Status update	1,320	420 chars.	
Wall post	1,323	1,000 chars.	
Event invite	1,237		Social graph structure.
Photo tag	1,103		Social graph structure.
Link	833		
Like	unknown		Social graph structure.
Note	unknown	65,536 chars.	Used for blog-style posts.

Table 2.1: *Facebook objects, their limitations and approximate frequency of creation [13]*

2.5.2 Connectedness

Since Encrypted Facebook’s broadcast encryption scheme has a transmission overhead proportional to the number of intended recipients, care must be taken to ensure the system works with large enough recipient groups. The number of recipients is bounded by the number of friends a user or equivalently by the degree of the user’s node in the social graph.

Empirical estimates for the average number of Facebook friends range from 130 to 170, with some evidence suggesting the distribution drops off sharply at around 250 [10] [7]. Marlow et al suggest that, regardless of the number of friends, communication is only ever between a small subset [2].

The Dunbar number is a theoretical cognitive limit to the number of peo-

ple a user can maintain relationships with and has been applied to social networks as well as face-to-face interactions. Exact estimates range from around 150 to 300 [xxx] [xxx], suggesting that the average degree of nodes within a social graph like Facebook's is unlikely to increase dramatically as Facebook expands further.

We this in mind we make it a requirement that the extension operates with recipient group sizes of up to 400.

2.5.3 Signal-to-noise ratio

Activity within the social graph causes notifications to be posted to feeds. Each user has a 'news feed' of nearby graph activity which is presented to them on logging in, and also a 'wall feed' of their own activity. Interaction with Facebook revolves around these feeds and it is vital that the signal-to-noise ratio is kept high. This is the proportion of useful content to non-useful content: typically spam, but in our case this could be transmission overheads as part of the broadcast encryption scheme, or even the encrypted content itself since a user without Encrypted Facebook installed will be unable to see the cleartext.

In order to permit incremental deployment any system must ensure that its users can coexist with non-users. We therefore make it a requirement to limit the extensions impact on the signal-to-noise ratio of feeds.

2.5.4 Graph API

Facebook does provide a JavaScript SDK for interfacing with the Facebook platform, however it is poorly documented and doesn't allow uploading images - since most JavaScript applications are designed to run inside the browser sandbox without local filesystem access. Instead we use the Facebook Graph API directly.

Objects can be read simply by making GET requests to `https://graph.facebook.com/ID` and parsing the return result as a JSON object. For none public objects we will also require an access token. Facebook uses the OAuth 2.0 protocol. An access token can be obtained by handling the page redirect³ after the following GET request:

³To ensure authentication can only occur in client side code the access token is passed in a URI fragment.

Listing 2.1: Authentication GET request

```
https://www.facebook.com/dialog/oauth?  
client_id=APP_ID&  
redirect_uri=REDIRECT_URL&  
scope=SCOPE_VAR1 , SCOPE_VAR2 , SCOPE_VAR3&  
response_type=token
```

Content can also be published in a similar way. An example request might look like:

Listing 2.2: Publishing POST request

```
https://graph.facebook.com/ID/feed?  
access_token=ACCESS_TOKEN&  
message=MESSAGE
```

There are several caveats:

- When publishing images, although the operation is supported, getting a correct handle to the image is difficult due to JavaScript's poor support for working with local files. The workaround requires creating an invisible form on the current page with a file input element and extracting the file handle from there.
- Images have to be published to an album. Facebook currently uses two types of album ID, one which appears within web pages and one which can be used for publishing through the Graph API. An additional API query is required before uploading to translate from one format to the other.
- In certain cases, though publishing through the Graph API is possible, it is more convenient to programmatically manipulate form controls. An example is when submitting a comment and triggering the click handler for the submit button.
- Modifying the "About Me" section of a user's profile is unsupported entirely. The workaround requires creating an invisible iframe on the current page and manipulating a form on the Facebook site within.

2.6 Storing data in images

As well as being one of the most frequently used types of content, images have been highlighted as a prime privacy concern on social networks [fb-images]. None of the existing work supports image encryption, though flyByNight mentions it as a possible extension.

The main problem is that, regardless of the input format, Facebook encodes all uploaded images using lossy JPEG.⁴ This means we require some form of JPEG-immune coding to store the binary output of encryption in an image, such that even after undergoing compression we can exactly recover the original bytes.

We begin by describing Facebook's JPEG compression process and evaluating naive attempts at encoding data in images, before proposing two more advanced approaches.

2.6.1 JPEG compression process

Information is lost at several stages:

1. Colour images are subject to a lossy colour space transform from RGB to YCrCb.
2. The chrominance components Cr and Cb are subsampled at a rate half that of the luminance channel.
3. The discrete cosine transform is applied to each 8x8 block using floating point arithmetic.
4. DCT coefficients are quantised according to values in a quantisation matrix.

Note that chrominance subsampling means that colour images only provide a 50% increase in capacity over a grayscale image of the same resolution. For simplicity we will only consider grayscale images. This leaves quantisation as the principle step at which information loss occurs.

Table 2.2 displays the quantisation matrix used for a grayscale JPEG image downloaded from Facebook. Using this and several other compression

⁴Even if a file is already in the output format the compression process is repeated and information is lost.

Quantisation Matrix							
5	3	3	5	7	12	15	18
4	4	4	6	8	17	18	17
4	4	5	7	12	17	21	17
4	5	7	9	15	26	24	19
5	7	11	17	20	33	31	23
7	11	17	19	24	31	34	28
15	19	23	26	31	36	36	30
22	28	29	29	34	30	31	30

Table 2.2: Quantisation matrix used by Facebook for luminance channel.

parameters our best guess is that Facebook is using the libjpeg library for compression, with a quality factor setting of 85.⁵

2.6.2 Naive data insertion

We encode data in an image and compress/decompress at quality factor 85, using libjpeg. We then examine the Hamming distance between the output and the original data and compute the bit error rate.

For each DCT coefficient with corresponding quantisation coefficient c we know that:

- $\lceil \log_2 n \rceil$ low order bits will be lost during quantisation.
- Setting higher order bits is undesirable due to clipping. Large magnitudes are more likely to produce values which fall outside the 0-255 interval when performing the inverse DCT.

Figure XXX graphs the error rates for naive insertion in to RGB values and insertion into DCT coefficients using a bitmask. Figure XXX also shows the per image capacity calculated by modelling the compression process as a binary symmetric channel (details for this are given in Evaluation section XXX).

⁵Based on the output of the JPEG Snoop application for Windows.

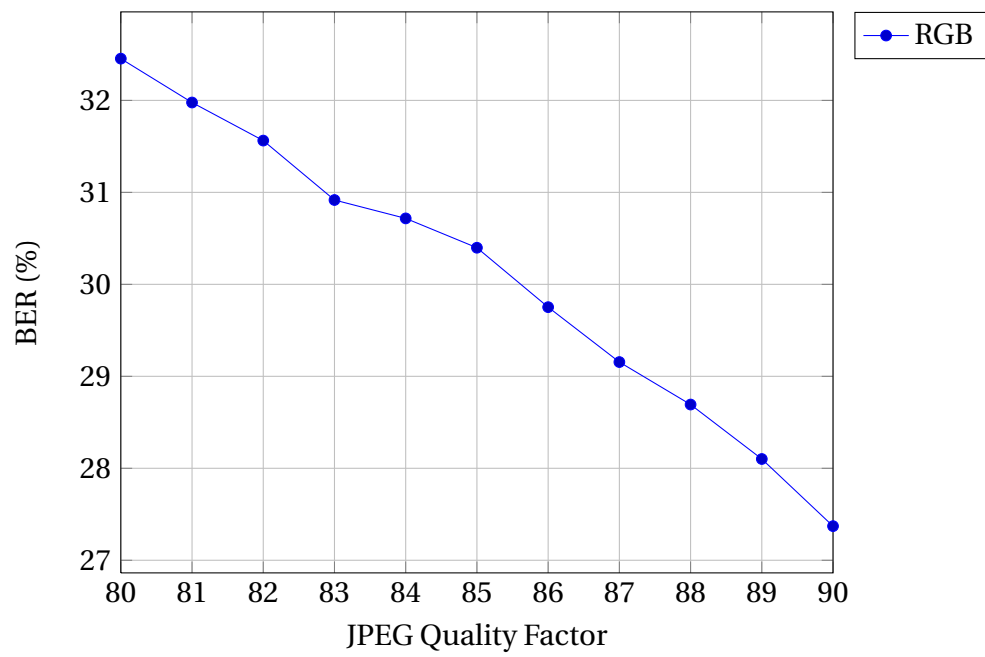


Figure 2.3: *Bit error rate for varying quality factors.*




Figure 2.4: *Per-image channel capacity (measured in KiB/image) for varying quality factors.*

2.6.3 Advanced data insertion

We map binary data on to appropriate length gray codes to ensure that only single bit errors occur from erroneously an adjacent codeword.

- JPEG DCT compression selectively quantises high frequency components since these are less perceptually salient. Wavelet transforms allow us to embed data in the low frequency sub-band of the carrier signal and can be performed reversibly (i.e. losslessly) by using an integer lifting scheme. Xu et al demonstrate that [16] data encoded in the low-frequency approximation coefficients can survive JPEG decompression when combined with an error correction scheme.
- What will from now on be referred to as a n -bit scaling method. We map the n -bit input space on to the 8-bit pixel space by scaling and shifting the input such that 0 corresponds to 0, and 2^n corresponds to 256. The inverse process amounts to outputting which interval the data lies in.

Clearly both these schemes are sub-optimal and their exact properties (compression time and error rates) are unknown. We make it a requirement that the project should take a modular approach to conduit image implementations to ensure that a) both of the proposed schemes can be implemented simultaneously and their performances compared and b) to aid future development of a more optimal solution.

2.7 Further security considerations

2.7.1 Threat model

We describe an attack centric threat model based on the methodology of [XXX]

2.7.2 Underlying encryption schemes

We use AES and RSA as both schemes are approved by NIST standards [nist] and widely available via the Botan cryptography library. Ideally however, we would have liked to use a scheme based on elliptic curve cryptography since public key sizes are much smaller for the same amount of security than finite field or integer factorisation. This is important since the block size of the cipher depends on the public key, and this in turn

determines the size of the session key once it has been encrypted. In our case, for example, a 256-bit AES session key requires 2048-bits of transmission overhead per recipient. Unfortunately ECC is less common due to patent concerns [XXX]. Again, we make it a requirement that the project be modular enough to allow later insertion of more optimal scheme.

2.7.3 Key management

- Key management and size (NIST recommendations).
- Message key and IVs, don't reuse. Ensure good source of entropy.
- Private key policy. Find a good reference, but basically we just mimic SSH and the like.
- Public key policy. Good idea to warn the user of the risks when they add public keys, check SSL is enabled etc.

2.8 Testing plan

- What kinds of testing will I use?
 - Unit testing , anything else??
 - Cognitive walkthrough - does this count as usability testing?
 - Security testing, since potential for exploits and project is security based - important enough to warrant its own section.
- How can I make these tests possible? Test bed or framework that needs to be in place?
 - Need a method of simulating the Facebook JPEG compression process. Use libjpeg since it most closely matches the compression signature (table of compression signatures). Show coefficient table.
 - Need a BER (bit error rate) calculator. Again coded as a C function.
 - FireBug and FireUnit for unit testing and profiling JavaScript functions.
 - gprof for profiling C/C++ functions.

2.9 Security testing

Loosely based on methodology here <http://mtc-m18.sid.inpe.br/col/sid.inpe.br/ePrint%4080/2006/12.20.12.15/doc/v1.pdf>. Must compromise since full security audit beyond scope of project. Look only at text retrieval process and public key management. We ignore images, and general attacks (e.g. setting up a spoof Facebook site). We also ignore threats that would be present ANYWAY e.g. if you haven't got SSL on. As an extension expand threat model.

- Threat analysis. $\text{Threat} = \text{Agent} \times \text{Mechanism} \times \text{Asset}$.
 - Facebook user creates a tag, which when decryption is attempted, causes denial of service (by locking up resources).
 - Facebook user creates a tag which when decrypted injects script in to page, gains control of users browser, can execute arbitrary scripts within the Facebook domain (XSS) gains access to Facebook cookies.
 - Facebook user exploits UTF8 encoder/decoder to smuggle illegal characters past sanitization, gains control of users browser, can execute arbitrary scripts within the Facebook domain (XSS) gains access to Facebook cookies.
 - Facebook user injects text which is run by JavaScripts eval() function, can execute arbitrary JavaScript outside the sandbox. Very Very bad!
 - Facebook user creates public key which, when parsed, creates a malicious file on the users local system.
- Risk analysis. $\text{Risk} = (\text{Vulnerability} \times \text{Threat} \times \text{Impact}) / \text{Security Measures}$.
 - Highest impact is running code outside the sandbox. True it maybe unlikely so long as we aren't stupid, but still. Basically we ban use of the eval function except for when we need it (retrieving JSON objects) then we replace it by a secureEval() which only allows valid Facebook object things.
 - Access to Facebook cookies can impact our security guarantees (since they could then change the public key). Also vulnerability is high. Thus we take time to sanitize before we inject into the browser.

- Denial of service is low impact, but high vulnerability since the user need not do anything to initiate the decoding process other than browsing to a site with a malicious post. So, test UTF8 decoder a lot, ensure that UTF8 decode, FEC decode, decrypt, all fail gracefully. Not image decode since out of scope, as mentioned above.
- Public keys we can limit to Base64 characters of a certain length. Done.
- Test plan elaboration. From the above we want:
 - Testing of secureEval. Override or otherwise ban eval().
 - Testing of text sanitiser.
 - Testing of UTF8 de/codec. Complicated given the large range of i/o.
 - Testing of public key downloader.

2.10 Professional practice stuff

- Software development methodology. Iterative prototyping. Work plan spells out which prototypes with what functionality should be completed when.
- Coding conventions, const correctness etc.
- Version/source control. Git and project locker.
- Performance bounds. Of what???

2.11 Requirements analysis

- Encryption should be available on the most commonly used tasks (apart from those otherwise ruled out in section XXX). The user must therefore be able to broadcast-encrypt, submit, retrieve and decipher the following objects.
 - Status updates
 - Wall posts
 - Comments

- Messages
- Images

Specifically, encryption should ensure confidentiality of data with at least 128 bits of security.

- All requirements should be met with recipient groups of size up to 400, which is a reasonable number - refer to discussion.
- Should be unobtrusive (refer to introduction) i.e. must not negatively affect browsing/Facebook experience of users. From this we derive the following:
 - Should try not to introduce any security holes. Up to a point, given scope of project. We have already declared a threat model and testing strategy etc.
 - Retrieval and submission times should be within acceptable limits. Define acceptable as <http://www.useit.com/papers/responsetime.html>.
 - Must not confine users to one computer. Should be portable. Securely transporting private keys is up to the user however.
- User activity should not negatively affect the activity of non-users (because of XXX refer to rest of preparation). We know there has to be some increase due to, for example, broadcast encryption overhead and status update's tiny length. Lets say maximum of twice number of objects generated compared to a normal user for the same activity.
- There are uncertainties and/or tradoffs associated with certain approaches to encryption and encoding data in images (and to a lesser extent error correction). It is also clear that in some cases the optimal approach is well beyond the scope of this project. Therefore, it is highly important that we adopt a modular structure that facilitates switching between differing schemes and permits future extension. This need not extend to simultaneously supporting different schemes - this would introduce much redundant complexity.

3 Implementation

3.1 Extension structure

We describe the overall structure of the extension and the C++ library. We also describe how the extension is loaded and initialised.

3.1.1 Overview

Aside from some boilerplate Firefox extension code and the JavaScript Stego! library (see section XXX) the main body of the application is made up of five components, shown on Figure 3.1:

Toolbar XUL

The toolbar XUL defines the toolbar interface using the Mozilla XML User Interface Language (XUL). The component is loaded as part of the browser chrome when Firefox is started and is responsible for loading the JavaScript components.

Page interception

pagecept contains the HTML parser for extracting prospective decryption targets and inserting UI controls. Actual target processors and control event handlers are defined in efb. This layer of abstraction means updates due to changes in the Facebook web site are isolated to this component. Component re-use (e.g. in an extension for another browser) is also facilitated.

Main extension component

efb defines the handlers for the toolbar and integrated UI controls. It contains handlers for decryption events and contains the plain-text cache data structures. It also contains callback handlers for asynchronous faceapi function calls. During the login process efb attaches the pagecept HTML parser to page loading events and

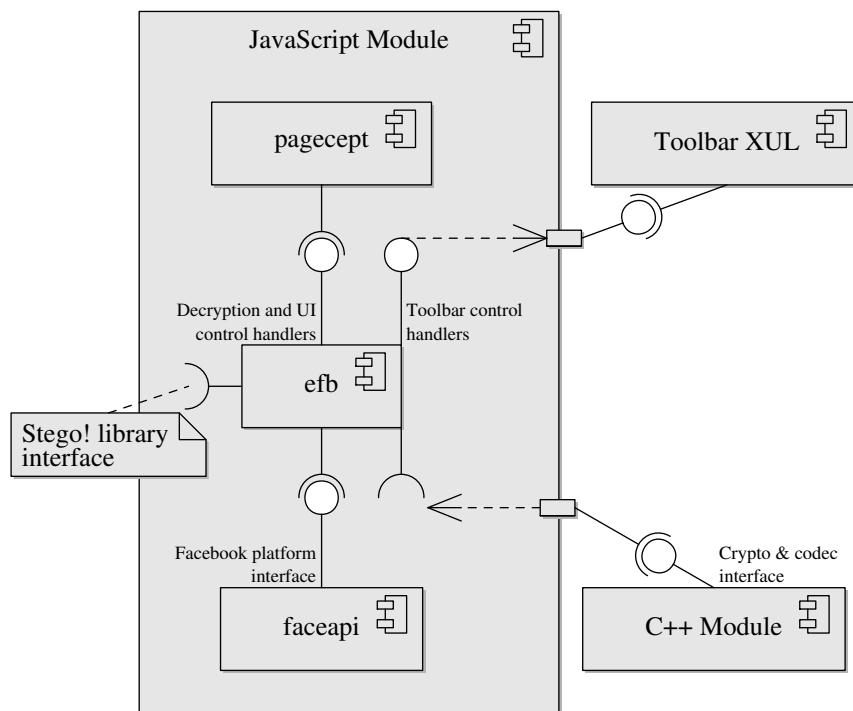


Figure 3.1: UML component diagram for the extension.

initialises the C++ module, binding the native methods to JavaScript wrapper functions.

Facebook API layer

`faceapi` is a layer of abstraction between `efb` and the Facebook platform. `faceapi` contains code for Graph API read/write queries as well as for the workaround solutions detailed in section XXX.

C++ Module

Primarily contains codec algorithms and cryptographic functions.

3.1.2 C++ module structure

The C++ module contains a library instance which implements the `IeFBLib` interface. The exposed behaviors of this library are wrapped appropriately so they may be called from the JavaScript module.

The library itself utilises four polymorphic sub-components:

- `ICrypto` contains cryptographic algorithms.

- `IFec` contains error correction algorithms.
- `IStringCodec` contains a UTF-8 encoder/decoder.
- `IConduitImage` contains JPEG-immune image coding algorithms.

This design facilitates future extension and possible run time composition of different sub-components - though currently the concrete implementations are chosen at design time.¹ The first three components are instantiated upon initialisation of the module. The last (`IConduitImage`) is generated whenever an image is encoded or decoded. Since C++ does not natively define interfaces we use an abstract base class with all pure virtual methods and a virtual destructor to disable polymorphic destruction [1].

The library is built around the abstract factory pattern described in [6]. This allows us to encapsulate groups of complementary sub-components since some interdependence exists between sub-components.² Figure 3.2 outlines the pattern structure with an example concrete subclass `HaarWTConduitImage`.

3.1.3 Initialisation

Firefox loads the toolbar XUL as part of the chrome when the browser is started. The XUL then loads the javascript components which define handlers for the toolbar controls. The extension is initialised when either the start button or generate identity button is clicked. `efb` harvests the Facebook ID from a browser cookie (assuming the user is logged in to Facebook) and uses it to define the working directory. The C++ library will then be instantiated and the `pacecept` handler attached to page load events.

¹Since only one set of components currently provides a feasible solution, see evaluation section XXX

²For example, the minimum size of the encryption header can't exceed the maximum capacity of the conduit image.

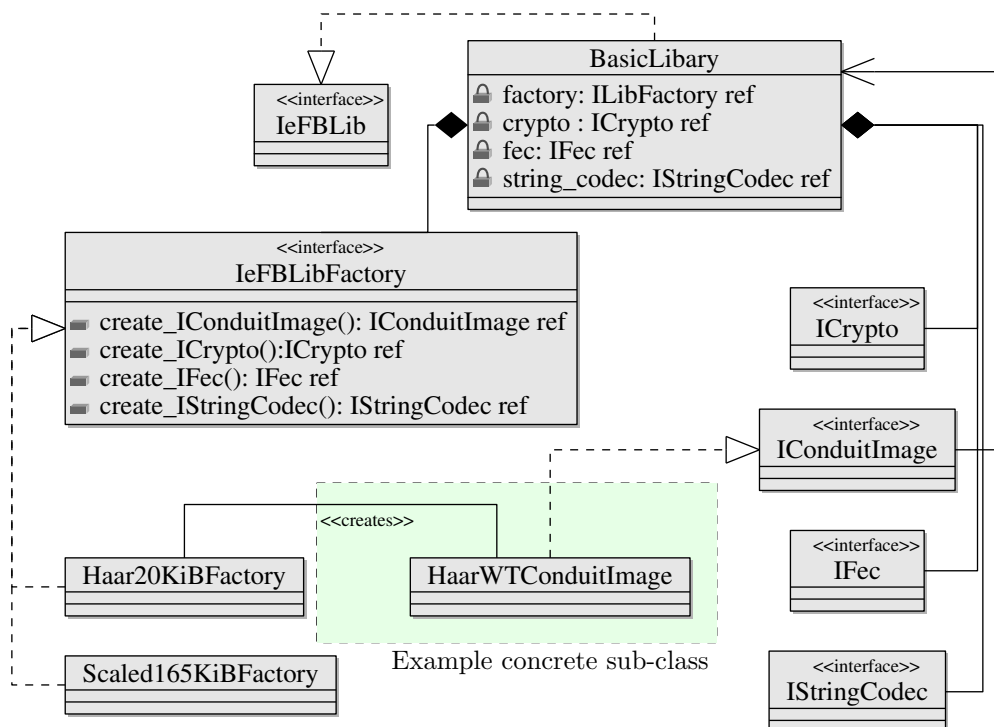


Figure 3.2: *UML class diagrams for the library and its sub-components.*

3.2 Integrating with the Facebook UI

The `pagecept` module checks if web pages are from the Facebook domain. If they are, it parses the HTML and attempts to insert additional UI controls and also retrieve and process any ciphered content automatically.

3.2.1 Inserting submission controls

There are four types of submission control, each of which may appear in multiple places within a single page. These are for general posts (status updates and wall posts) comments, private messages and image uploads. Each control is associated with an input field. Finding control-field pairs involves several complex regex queries which we will not discuss in detail.

Once a pair has been identified we generate an alternative encrypted submission control and place it beside the normal one. A handler is generated associated with the input field. For images this process is slightly

different - a check box is added and the handler to the normal control modified.

The submission handler causes a friend selector window to appear, loaded with any friends whose public keys are stored on disk. Elements within the control are populated with the user's names and profile pictures by performing queries through `faceapi`. Optionally, Encrypted Facebook will check if local public keys are up to date with online keys. The user is given the option of updating an out of date key, and is informed of the tradeoff between vulnerability to middleperson attacks and potential non-availability.

On submission, the selected list of friends and the input from the input fields are gathered and processed by `efb`. The submission process is detailed in section XXX.

3.2.2 Retrieving content

Regexes are used to locate and filter possible decryption targets. For text these will be enclosed by special start and end sequences. Images can be identified by their filename as Graph API objects. In either case we can not be certain and so filtering is a "best effort" approach. Malicious user could easily create fake text tags, and most Graph API images won't be encrypted. The decryption process is designed to fail gracefully as early as possible if this turns out to be the case. Currently, since all images are 720x720 pixels this predicate is used as a filter. The case for uploading variable sized images and its effect on filtering are discussed in section XXX.

Once a list of target Facebook IDs has been generated and filtered, it is processed. Each ID is checked to see if it has an entry in the cache. If not, an entry is created and a `XmlHttpRequest` triggered through `efb`. A handler is attached to the request so that on completion, the cache can be updated appropriately. If an entry exists then several actions may be appropriate. If a valid plaintext exists in the cache this is used. The entry may also be marked as in progress in which case a loading message is substituted. If a previous attempt failed then the target can be ignored.

3.3 Text submission

We consider the process by which a plaintext message, given a list of recipients, is transformed to a steganography encoded tag which points to

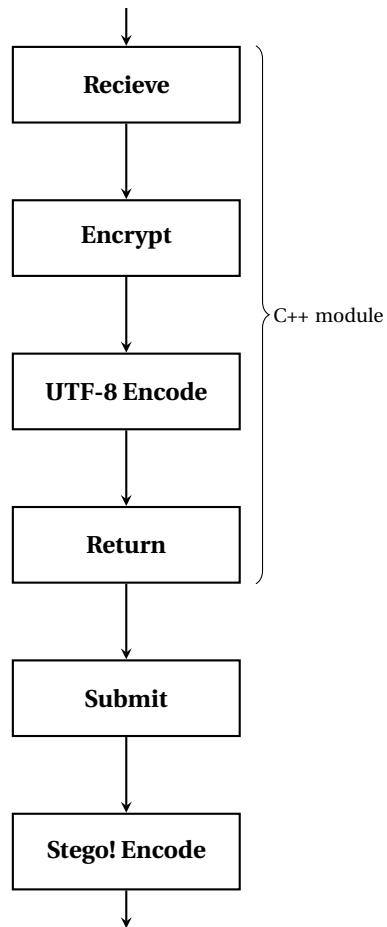


Figure 3.3: *Encoding process for submitting text.*

the ciphered message stored online as a note. This tag can then be used in place of the plaintext.

The string is passed to the C++ library, encrypted then encoded in a UTF-8 based format suitable for Facebook. The result is returned to the JavaScript module, submitted and a tag generated from the resultant object ID using steganography. We describe each step of this process in detail.




Figure 3.4: UML class diagrams for the cryptography component.

Description	Size (bytes)
Length tag	2
Initialisation vector	16
Facebook ID	8
Session key	<pub-key size>
⋮	
Facebook ID	8
Session key	<pub-key size>

Table 3.1: Structure of the encryption header.

3.3.1 Encryption

Currently the only implementation for cryptographic functions is based on the Botan library using RSA and AES. The Botan library is encapsulated in a class with template parameters (N , M) which determine the length (in bytes) of the AES session key and RSA public key, respectively. The class is designed so that a class with certain key sizes can be defined simply by specifying these parameters.

The input string is converted to a byte vector containing enough free space at the beginning for the encryption header (the size of the crypto header can be calculated in advance based on the recipient list so that encryption can be performed in place) along with a recipient list of Facebook IDs. The output is the ciphered message with the encryption header

prepended.

Table 3.1 describes the format of the crypto header generated as part of the broadcast encryption scheme. Note that the public key size determines the cipher block size and therefore the storage requirements for the encrypted session key - regardless of the actual session key length itself.

The Botan SecureVector data structure is used to intermediately store all cryptographic keys, preventing key material being swapped to disk. A random IV and session key is generated for every message using Botan, which is supposedly reasonably random [XXX]. After encryption, all seeds, key material and IVs are disposed of securely.

3.3.2 String coding

The input is a byte vector of random bytes from the encryption stage. Each 16-bit (2 bytes) code is mapped on to a valid UTF-8 character - a variable length sequence of 1 to 5 bytes. Odd numbered input is padded and a otherwise unused character sequence prepended to indicate this. The mapping is based on the mapping from Unicode code points to UTF-8 chars, with two distinctions.

- Each 16-bit input is shifted by an offset of 0xB0 before being mapped to a character. This avoids problem symbol characters which will be escaped by the Facebook sanitization process (< and > for example).
- Unicode code points XXX-XXX are surrogate pair characters and are illegal if used in isolation. Inputs which map to these characters (after being offset) are bit-shifted left by one place.

Note that this means some of the resulting code points are outside the BMP (Basic Multilingual Plane) but are still well defined as UTF-8 characters and supported by Facebook.

After adding a null terminal the final string can be returned to the JavaScript calling function.

3.3.3 Submission as a note

The final string is submitted as a note to Facebook via `faceapi` passing the relevant handlers from `efb`. On completion the Facebook Graph API object ID is parsed from the `XmlHttpRequest` response and encoded using the steganography functions provided by the JavaScript `Stego!` library. Start and end tags are added and the final text is ready to be used in place of the cleartext, as described in section XXX.




Figure 3.5: *Encoding process for submitting an image.*

3.4 Image submission

We now describe the process by which an image, stored locally, is encrypted and encoded in a temporary image file ready to be uploaded. The C++ library is passed the input and output file paths and returns 0 on success.

Initially the image data is loaded from disk as a byte vector (leaving room for the encryption header) and encrypted exactly as described in section XXX. Error correcting codes are then added. Finally, a conduit image object is created, written to, and saved to disk. We describe the last two stages in detail.

3.4.1 Error correction

Currently the only implementations are based on the Shifra library using Reed Solomon codes. The bulk of the Shifra library is encapsulated in an

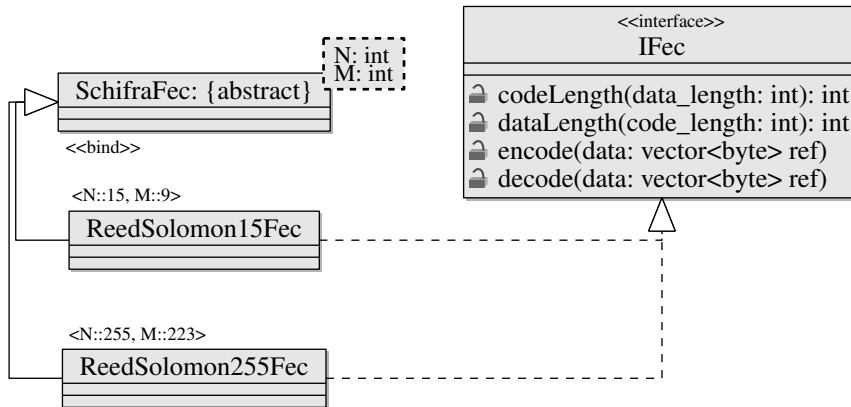


Figure 3.6: UML class diagrams for the forward error correction library component.

abstract base class and two template specialisation subclasses implementing codes rates of (15,9) and (255,223) (see Figure 3.6).

The FEC algorithms have a fixed block size. In order to avoid excess padding bytes or addition length tags, we use a scheme inspired by cipher-text stealing used in cryptography [XXX]. The FEC codes are appended to the data bytes so that, provided there are enough blocks, any last partial block will be padded out automatically.

3.4.2 Conduit image class hierarchy

Images are encoded by generating an instance of IConduitImage through the abstract factory, writing data to it, then saving out to disk. The CImg class (from the CImg library) is used as a base class since it supports opening and saving various image formats, manipulating pixels and colour space transforms.

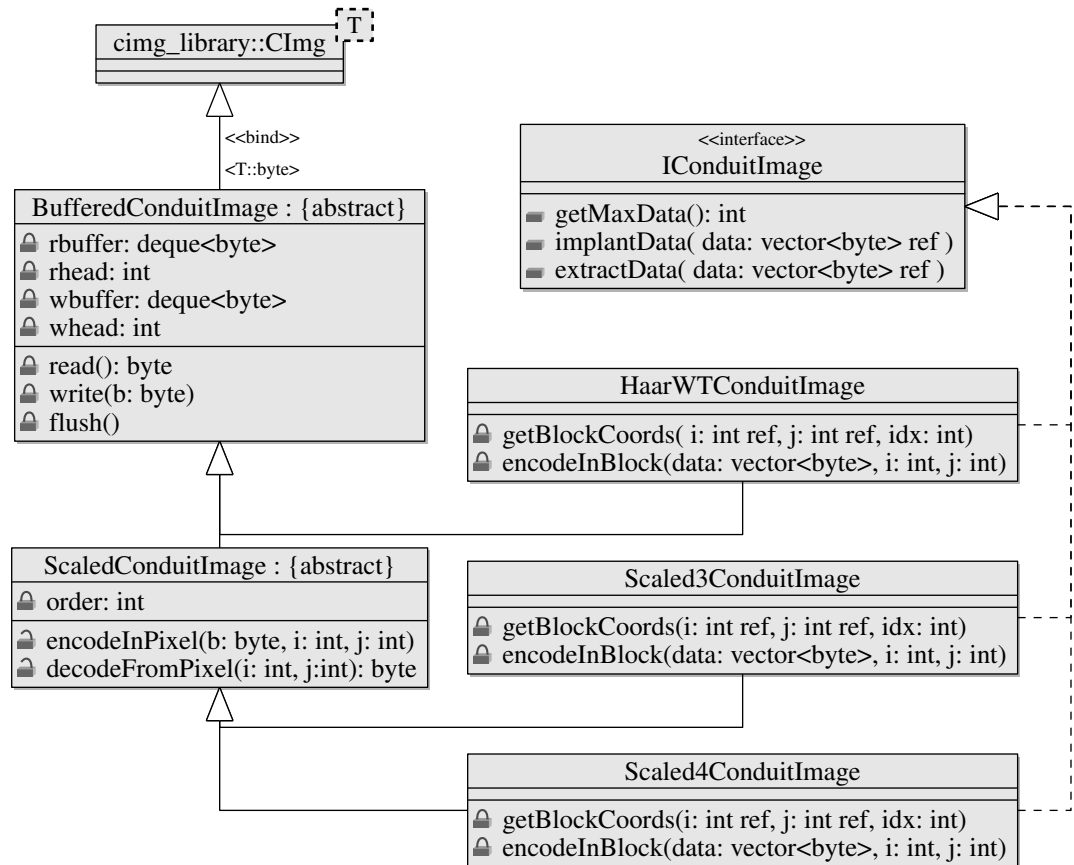


Figure 3.7: UML class diagrams for the conduit image implementation.

3.4.3 Read/write buffering

All current implementations are derived from the abstract class `BufferedConduitImage` which implements read and write buffers. This allows us to group single byte read/write requests together, which is necessary for certain implementations. The Haar wavelet method, for example, stores 3-bytes at a time in 8x8 pixel blocks.

The variables `rhead` and `whead` determine the current position of the read and write heads, or equivalently the number of bytes written or read since creation. `block_size` is set upon instantiation through the constructor and is the smallest integer number of bytes that can be read/written at once.

Any subclass of `BufferedConduitImage` must provide a function `getBlockCoords`. This maps the position of the read/write heads to block coordinates within the image. Subclasses must also implement `encodeInBlock` and `decodeFromBlock` for writing `block_size` bytes to and from a block given the block coordinates.

This class also contains gray code translation functions, since they are used by all descendant classes and their definitions are too small to justify a class of their own.

Class	Dimensions (pixels)	Block size (bytes)	Grey codes (bits)
Haar WT	8×8	3	6
3-bit Scaling	3×2	3	3
4-bit Scaling	2×1	1	4

Table 3.2: Comparison of blocks for each concrete subclass

3.4.4 Haar wavelet transform

The `HaarWTConduitImage` uses blocks of 8x8 pixels, with a block size of 3-bytes. Two passes of the 2D Haar wavelet transform are performed on a single block. 6 bits are written to the high order bits of each of the four 8-bit approximation coefficients. The low order bits are masked off based on experimental results, and as suggested in [XXX]. The inverse transform is then performed to output grayscale pixel values.

The `CImg` class does contain Haar transforms but these are not suitable. We require an integer lifting scheme (described in [XXX]) to ensure that the

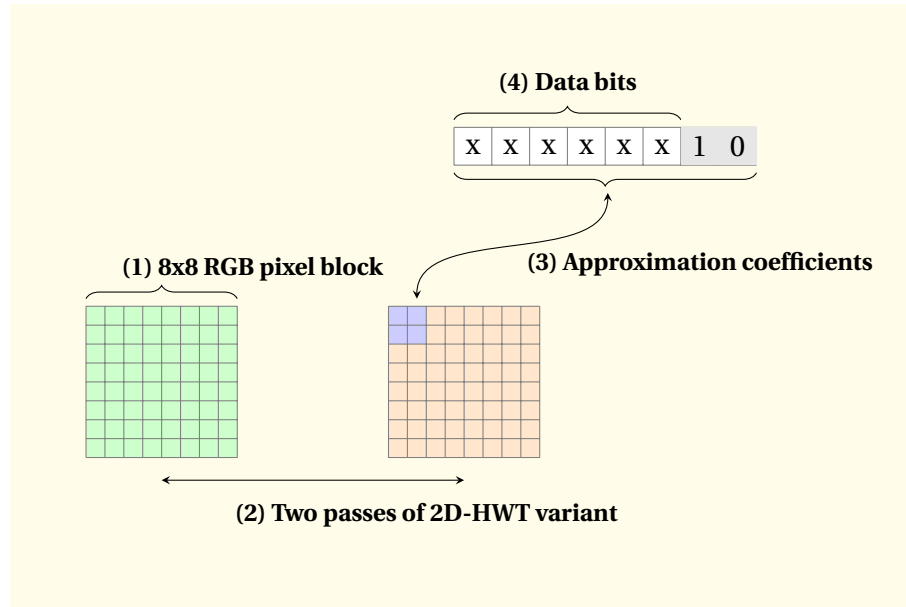


Figure 3.8: Outline of the encoding process

transform is reversible (without information loss). In the 1-dimensional case, for a pair of approximation and difference coefficients (c_a, c_d), we calculate the output value pair (x, y):

$$\begin{aligned} x &= c_a + \lfloor \frac{c_d+1}{2} \rfloor \\ y &= x - c_d \end{aligned} \quad (3.1)$$

Iteration of the above over pairs in both vertical and horizontal axis can be used to perform the full 2D HWT and its inversion losslessly.

Changing the approximation coefficients can also lead to capping of values as they are mapped back to greyscale space, outside the 0-255 range. We therefore selectively discard high frequency information (during the inverse transform) from the difference coefficient c_d whenever capping occurs - leaving the approximation coefficients intact and the output within range. See appendix XXX for full details of the exact HWT variant used.

3.4.5 N-bit scaling

The base class `ScaledConduitImage` contains functions to code n-bits of data to and from a single pixel. The value of n is set on instantiation.

Two subclasses are specified for $n = 3$ and $n = 4$. For $n = 3$ we use 3x2 blocks of pixels with a block size of 3-bytes. For $n = 4$ we use blocks of 2 pixels with a block size of 1-byte.

The scaling process works so that the intervals for each data point are of equal width, except for the intervals at either end which are $\frac{1}{2}$ length. This is because extreme values (0 or 255) can only be either decreased or increased due to compression artifacts, not both. An input pixel value of 255 might result in 254 or 253 after compression, but never 0 or 1.

3.5 Testing

- Unit
- Regression
- Black box
- White box
- Integration
- Security/penetration testing?

Security tests

- Use of the `eval()` and `secureEval()` functions
 - test `secureEval()`. Should only decode JSON objects. Should only do so from Facebook API requests.
- Insertion of text in to page. Easy since we can use JavaScript and RegExp.
 - We allow all uppercase lowercase letters and numerals. Also allow `.,?!()`. That's it, better safe than sorry. Means no linking to malicious pages. Fully test all boundary cases etc etc.
- UTF-decoder. Slightly harder since have to look at bytes not characters. Using the following rules we conformance test, test all boundaries etc etc. Put list of test inputs in appendix.
 - We accept any valid, non-overlong, UTF8 byte sequences, max length 4-bytes, with scalar value:
 - * 0xB0 - 0xD7FF

- * 0xE000 - 0x100AF
- * 0x1B000 - 0x1BFFE (would-be surrogate pairs)
- * 0x10F0000 (indicates a padding byte was added, only one allowed per decode)
- We therefore must throw an exception whenever a valid UTF8 byte sequence is presented with scalar value:
 - * 0x0 - 0xAF (out of range)
 - * 0xD800 - 0xDFFF (surrogate pair characters)
 - * 0x100B0 - 0x1AFFF (out of range)
 - * 0x1BFFF - 0x10FFFFFF (out of range)
 - * 0x10F001 - 0x1FFFFFF (out of range)
- We also throw an exception for valid UTF8 sequences when:
 - * They have an overlong form i.e. the same scalar value can be represented using a shorter byte sequence.
 - * They have scalar value 0x10F0000 (padding character) but this has already been seen during decoding.
 - * They have scalar value 0x10F0000 (padding character) but the final decoded byte sequence (before padding removal) has length less than 2.
 - * The final decoded byte sequence has length less than 1.
 - * They are longer than 4-bytes.
- Naturally we reject any (invalid) UTF8 byte sequences with:
 - * Unexpected continuation bytes when we expect a start character.
 - * A start character which is not followed by the appropriate amount of valid continuation bytes - including start characters right at the end of a sequence.
- Public key downloader. Simply limit size, don't use exact size since other implementation might use different key sizes.